

Simulating Secure DNN Inference on Fully Homomorphic Encryption

Woosung Song

Seoul National University, Dept. of Computer Science and Engineering

December 10, 2021

Simulating Secure DNN Inference on FHE

- Simulating DNN inference as if it is on a secure FHE environment.
- Switching HE-incompatible ML components into approximate polynomials.
 - ▶ E.g., Conditional branches.
- Optimize and fine-tune the model to reach a higher accuracy.
- Compare the tradeoffs between accuracies and computational costs.

Homomorphic Encryption

Homomorphism Between Plaintexts and Ciphertexts

- Can perform operations on the ciphertexts without decryption
- Homomorphism between plaintext and ciphertext spaces

Let $c_1 \in Enc_{pk}(p_1)$, $c_2 = Enc_{pk}(p_2)$.

$$Dec_{sk}(c_1 \pm c_2) = p_1 \pm p_2$$

$$Dec_{sk}(c_1 \times c_2) = p_1 \times p_2$$

$$Dec_{sk}(f(c_1, c_2)) = f(p_1, p_2)$$

Fully Homomorphic Encryption

HEaaN (CKKS)

- Fully HE scheme for **floating-point approximate arithmetics**
- Supports $+$, \times , *rotate* operators for vectorized data packets (SIMD)

$$(a_1, a_2, \dots, a_n) + (b_1, b_2, \dots, b_n) = (a_1 + b_1, a_2 + b_2, \dots, a_n + b_n)$$

$$(a_1, a_2, \dots, a_n) \times (b_1, b_2, \dots, b_n) = (a_1 \times b_1, a_2 \times b_2, \dots, a_n \times b_n)$$

$$\text{rotate}((a_1, a_2, \dots, a_n), r) = (a_{r+1}, \dots, a_n, a_1, \dots, a_r)$$

$$\text{Dec}_{sk}(c_1 \pm c_2) = p_1 + p_2$$

$$\text{Dec}_{sk}(c_1 \times c_2) = p_1 \times p_2$$

$$\text{Dec}_{sk}(\text{rotate}(c_1, r)) = \text{rotate}(p_1, r)$$

Noise Level & Bootstrapping

Noise Level

- Ciphertext's noise(error) becomes significantly larger after a multiplication
- Ciphertext lost its information when the level goes 0

Bootstrapping

- Operator in HEaAN that **recovers the noise level**
- Takes very very long time...
 - ▶ This is why one should care to **reduce the cascaded multiplicative depths**
- Enables Fully HE: No need to decrypt ciphertexts with low noise levels

Let \mathcal{C}_l be a set of ciphertexts with l noise level.

$$Enc_{pk}(p) \in \mathcal{C}_{init}$$

(Initial (large) noise level)

$$\pm : \mathcal{C}_l \times \mathcal{C}_{l'} \rightarrow \mathcal{C}_{\min\{l, l'\}}$$

$$\underline{\times} : \mathcal{C}_l \times \mathcal{C}_{l'} \rightarrow \mathcal{C}_{\min\{l, l'\}-1}$$

$$\underline{rotate} : \mathcal{C}_l \times \mathbb{Z} \rightarrow \mathcal{C}_l$$

$$\underline{bootstrap} : \mathcal{C}_{low} \rightarrow \mathcal{C}_{high}$$

Secure DNN Inference on FHE

DNN Inference using FHE Scheme

- DNN inference from encrypted user data
- Using homomorphic operators
 - ▶ $+$, \times , *rotate*
- Popular DNN components on inference stage:
 - ▶ Matmul, Conv, Reshape, Sum ...
 - ▶ ReLU, MaxPool, Tanh, ...




Figure: Secure DNN inference on FHE scheme

Secure DNN Inference on FHE

No Conditional Instruction in HEaaN?

- Don't know any information from the ciphertext
 - ▶ How can one implement ' $ReLU(x) = \text{if } (x \geq 0) \text{ then } x \text{ else } 0$ '?

Approximate Polynomials

- Approximate polynomials to implement **the absolute function**
 - ▶ $|x| \approx a_0 + a_1x + a_2x^2 + \dots$
- Maximum degree of the approx polynomial determines noise level drop
 - ▶ **Tradeoff between multiplicative depth v.s. accuracy**

$$ReLU(x) = (x + |x|) \cdot 0.5$$

$$MaxPool(x_1, x_2) = \max(x_1, x_2) = 0.5 \cdot (x_1 + x_2 + |x_1 - x_2|)$$

Chebyshev's Approximation (1st Kind Expansion)

- Known to have low average error ($\text{mean}(|f - p|)$).
- Recursively generates polynomials.

$$T_0(x) = 1, \quad T_1(x) = x, \quad T_{n+1}(x) = 2xT_n(x) - T_{n-1}(x)$$
$$\sum_{k=1}^n T_i(x_k) T_j(x_k) = \begin{cases} 0 & \text{if } i \neq j \leq n \\ n & \text{if } i = j = 0 \\ n/2 & \text{if } 0 < i = j \leq n \end{cases} \quad \text{where } x_k = -\cos\left(\frac{k\pi}{n}\right)$$

- Compute approximation coefficients similar to Fourier transformation.

$$f(x) \approx \sum_{i=0}^n c_i T_i(x), \quad c_i = \frac{1}{d_i} \sum_{k=1}^n f(x_k) T_i(x_k), \quad d_i = \begin{cases} n & \text{if } i = 0 \\ n/2 & \text{if } i \neq 0 \end{cases}$$

Remez's Iterative Algorithm

Theorem (Chebyshev's Equioscillation theorem)

Among the approximate polynomial p of f s.t. $\deg(p) \leq n$, $\|f - p\|_\infty$ on a domain $[a, b]$ is minimized iff $\exists a \leq x_1 < x_2 < \dots < x_{n+2} \leq b$ s.t. $f(x_i) - p(x_i) = (-1)^i \cdot \|f - p\|_\infty$.

- Known to have minimal maximum error ($\max(|f - p|)$).
- Iteratively fetches x points and polynomial coefficients.
 - ▶ Randomly sample $n + 2$ ascending data points $(x_1, x_2, \dots, x_{n+2})$ on $[-1, 1]$.
 - ▶ Solve a linear system of $n + 2$ equations, with $n + 2$ unknowns c_0, \dots, c_n, E :

$$\left(\sum_{i=0}^n c_i x_k^i \right) + (-1)^k E = f(x_k)$$

- ▶ Update $(x_1, x_2, \dots, x_{n+2})$ to have local maximums on $\|f(x) - p(x)\|$, and repeat from the 2nd step.

Approximations of Absolute Function





Figure: Chebyshev's (L), Remez's approximation of absolute function $y = |x|$ (R)

- Chebyshev has low average error, Remez's has low maximum error.
- Higher degree polynomials are more accurate.

Importance of Input Range Issue

Recall that $\text{ReLU}(x) = (x + |x|) \cdot 0.5$

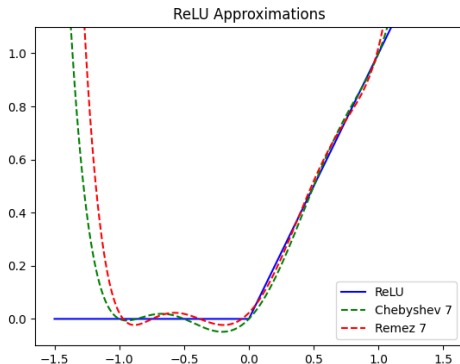


Figure: Approximations of ReLU. The function is highly inaccurate out of $[-1, 1]$ input range.

Experiment Environments

Target Datasets and Networks

- MNIST handwritten digit classification dataset.
- Simple convolutional classifier neural network.
 - ▶ Conv2d, Linear, ReLU and MaxPool2d in pytorch.
- Implemented with PyTorch for secure inference simulation.
 - ▶ Use approximation polynomials when inferencing.
 - ▶ Not necessary for training phase.




Figure: Classifier Network

Optimization & Fine-Tuning

ML Components and Losses

- **Activation decay**; to reduce input norms of approx. polynomials.

$$Loss(\theta, x) = Loss_{orig}(\theta, x) + \lambda_{act} \cdot \left(\frac{1}{N} \cdot \sum_{u \in \text{inputs of approx.}} ||u||_p \right)$$

- Negative slope ($u \in [0, 1)$) on ReLU (LeakyReLU)

$$LeakyReLU(x) = \begin{cases} x & \text{if } x \geq 0 \\ ux & \text{if } x < 0 \end{cases} = \left(x + \frac{1-u}{1+u} \cdot |x| \right) \cdot \frac{1+u}{2}$$

Fine-tuning

- Fine-tune with approximates
- Gradient descent training after switching HE-incompatible components into approximates

Experiment Pipelines

Common Training Pipelines

- 1 Train the network w/ small activation decay ($\lambda_{act,1} = 10^{-4}$) for 20 epochs
 - ▶ Learning rates: 0.1(10 epochs) \rightarrow 0.01(15 epochs) \rightarrow 0.001(20 epochs)
- 2 Train the network w/ **large activation decay** ($\lambda_{act,2} = 10^{-2}$) for 10 epochs
 - ▶ Learning rates: $10^{-3} \rightarrow 10^{-4} \rightarrow 10^{-5}$
- 3 **Fine-tune the network by replacing** HE-incompatible components into approximates for 10 epochs ($\lambda_{act,3} = 10^{-2}$).
 - ▶ Learning rates: $10^{-3} \rightarrow 10^{-4} \rightarrow 10^{-5}$

Control Variables

- Approximation methods (Chebyshev v.s. Remez) and degrees
- Using LeakyReLU instead of ReLU
- Using $L2$ or $L1$ -norm as activation decay
- Final activation decay parameter scale ($\lambda_{act,3}$)

Github: https://github.com/lego0901/fhe_secure_inference_simulation

Results - LeakyReLU on L1 Activation Decay




Figure: Affects of LeakyReLU. ReLU was more stable for $L1$ activation decay.

Results - LeakyReLU on L2 Activation Decay




Figure: Affects of LeakyReLU. LeakyReLU was generally better for L_2 activation decay.

Results - L1 or L2 Activation Decay Norms




Figure: Affects of L1 and L2 activation decay norms. L2 gave impressive performances for some cases, but L1 was more stable.

Results - Activation Decay Strength ($\lambda_{act,3}$)

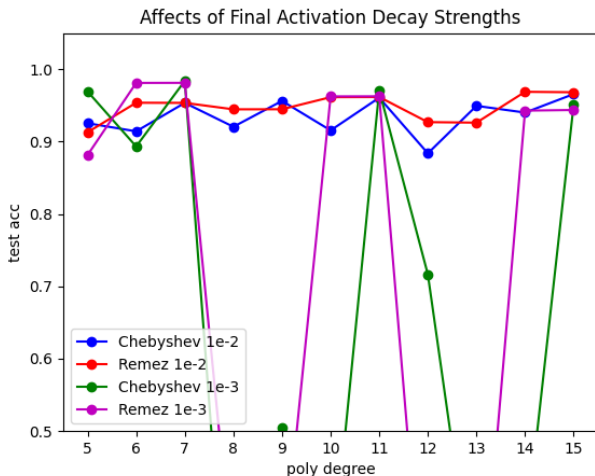


Figure: Affects of the last activation decay term ($\lambda_{act,3}$). 10^{-3} gave impressive performances for some cases, but 10^{-2} was more stable.

Results - Affects of The Fine-tune Phase

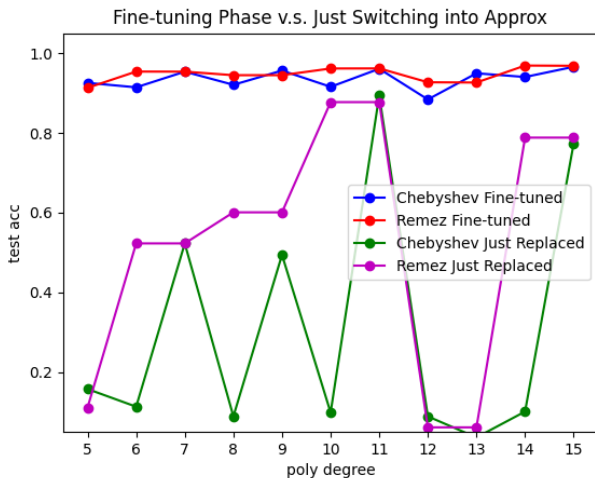


Figure: Affects of the fine-tuning phase. Just replacing ReLU and MaxPool2d into Approximates without fine-tuning gave poor performances.

Results - Affects of The Large Activation Decay Phase




Figure: Affects of the (2nd) large activation decay phase. It provided more stable results for smaller approximation polynomial degrees.

Results - Activation Magnitude Profiles




Figure: Activation magnitude profiles. The absolute value of each activation tensor element was investigated. Fine-tuning phase reduced the activation magnitudes in general.

Results - Top Accuracies Configs

Approximates	Act Decay	ReLU Layer	Accuracies
Chebyshev 7	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.71
Chebyshev 8	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.67
Chebyshev 6	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.63
Chebyshev 7	L1, $\lambda_{act,3} = 10^{-3}$	ReLU	98.39
Remez 7	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.31
Remez 6	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.31
Remez 10	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.21
Remez 11	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.21
Chebyshev 5	L2, $\lambda_{act,3} = 10^{-2}$	LeakyReLU	98.12
Remez 5	L1, $\lambda_{act,3} = 10^{-3}$	ReLU	98.12
Remez 6	L1, $\lambda_{act,3} = 10^{-3}$	ReLU	98.11

Table: Configurations with top accuracies

Results - Activation Magnitude Profiles

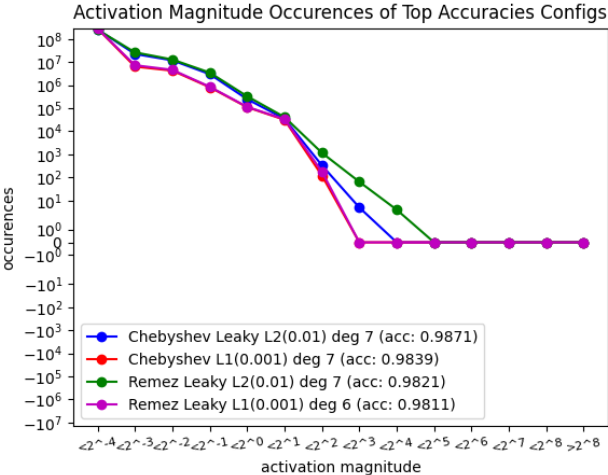


Figure: Activation magnitude profiles for remarkable configurations.

Conclusion

Summary

- Can switch HE-incompatible components into approximates.
- The accuracy drops were not critical in MNIST dataset. (98.71% in best)
- Fine-tuning phases were important to obtain better results.
 - ▶ With 'large activation decay' and 'fine-tuning' phases.
 - ▶ Fine-tuning stages reduced the magnitudes of activations, to fit into valid input ranges for the approximations; $[-1, 1]$.
- Low-degree polynomials approximations were enough. (Even better!)
 - ▶ About 6-7 degrees polynomials produced the best performances.
 - ▶ Better for all computational complexities. (resources, times, mult-depths)

Follow-up Studies

- Not 'simulating' a secure ML inference in FHE.
 - ▶ Need to implement `Matmul`, `Conv2d`, ...
- How to optimize these technique into more complex networks.
- How to make a secure training network? (not only the inference)

References I

- [1] Jung Hee Cheon et al. “Bootstrapping for Approximate Homomorphic Encryption”. In: *Advances in Cryptology – EUROCRYPT 2018*. Ed. by Jesper Buus Nielsen and Vincent Rijmen. Cham: Springer International Publishing, 2018, pp. 360–384. ISBN: 978-3-319-78381-9.
- [2] Jung Hee Cheon et al. “Homomorphic Encryption for Arithmetic of Approximate Numbers”. In: *Advances in Cryptology – ASIACRYPT 2017*. Ed. by Tsuyoshi Takagi and Thomas Peyrin. Cham: Springer International Publishing, 2017, pp. 409–437. ISBN: 978-3-319-70694-8.
- [3] Jung Hee Cheon et al. “Numerical Method for Comparison on Homomorphically Encrypted Numbers”. In: *Advances in Cryptology – ASIACRYPT 2019*. Ed. by Steven D. Galbraith and Shiho Moriai. Cham: Springer International Publishing, 2019, pp. 415–445. ISBN: 978-3-030-34621-8.
- [4] Thore Graepel, Kristin Lauter, and Michael Naehrig. “ML Confidential: Machine Learning on Encrypted Data”. In: *Information Security and Cryptology – ICISC 2012*. Ed. by Taekyoung Kwon, Mun-Kyu Lee, and Daesung Kwon. Berlin, Heidelberg: Springer Berlin Heidelberg, 2013, pp. 1–21. ISBN: 978-3-642-37682-5.
- [5] M.H. Mudde, Alef E. Sterk, and van der Schaft Arjan. “Chebyshev approximation”. In: (2017).